Relational Databases

aws training and certification

SQL Overview

aws training and certification

Data Definition Language (DDL)

Data Control Language (DCL)

Data Manipulation Language (DML)

Structured Query Language (SQL)

Transactional Databases

aws training and certification

- Relational Database
- Collects and manages transactional and operational data
- Examples include
  - Point-of-Sale (POS) systems
  - Registration systems
  - Auditing systems

Analytical Databases

- Relational Database
- Performs complex analytical queries
- Data is imported from other, often transactional, systems

Comparison of OLTP and OLAP Systems

aws training and certification

	Transactional	Analytical
Data source	Origin	Consumer
Purpose	Capture data	Analyze data
Workloads	INSERT, UPDATE, DELETE, short and fast queries	Batch jobs to import data, JOINs, complex queries
Database design	Highly normalized using many distinct tables to reduce duplication	Denormalized using fewer tables in star and snowflake schemas with some duplicated data
Database size	Depends on the amount of data but typically from MB to TB in size	Grows over time and typically ranges from TB to PB in size.

Data Warehousing Concepts

aws training and certification

Data and Insights

aws training and certification

Data is not information.

Data sources

Data transformations

Analytical data

What is a Data Warehouse?

A data warehouse is:

- A central repository of business data from disparate sources.
- A type of relational database that enables analysis of data.
- A collection of approved and trusted historical corporate data.

Data Warehousing Goals

- Provide access to corporate or organizational data.
- Ensure consistency and quality of the data.
- Enable analysis of data in different ways based on measurements that are defined by the business.
- Integrate with query, analysis, visualization, and reporting tools.

What Are Data Marts?

- A subset of a data warehouse
- Used to model a single subject or dimension, such as product or customer
- Often organized in a star or snowflake schema
- Fed from a data warehouse (top-down approach) or feeds a data warehouse (bottom-up approach)

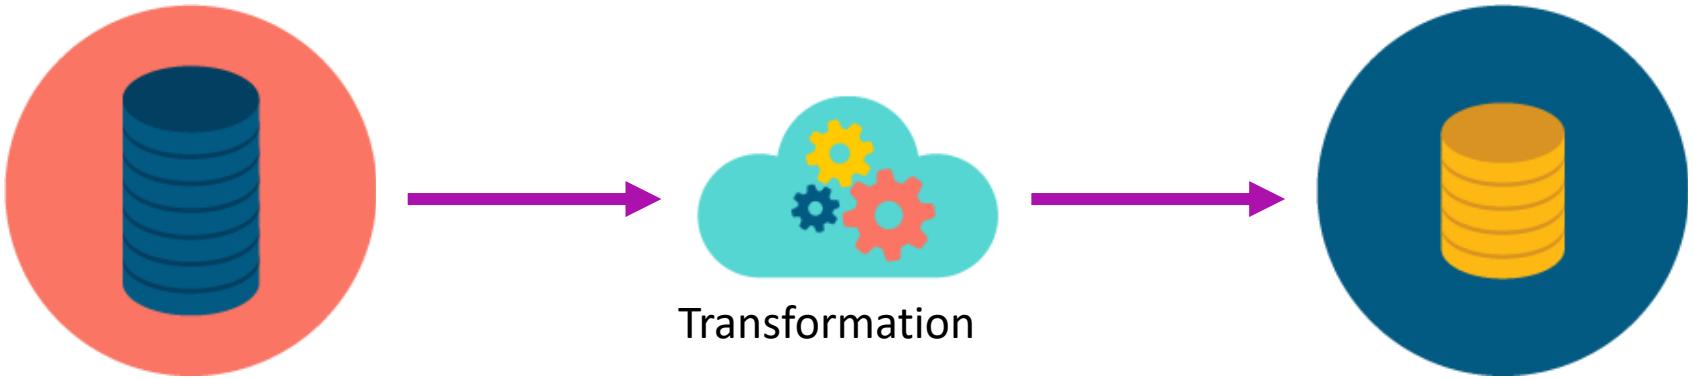

Dimensional Modeling: Star Schema

- Developed by data warehouse architect, Ralph Kimball.
- The center is a fact table and the points are dimension tables.
- Many business intelligence (BI) applications support the star schema.

Star Schema Example

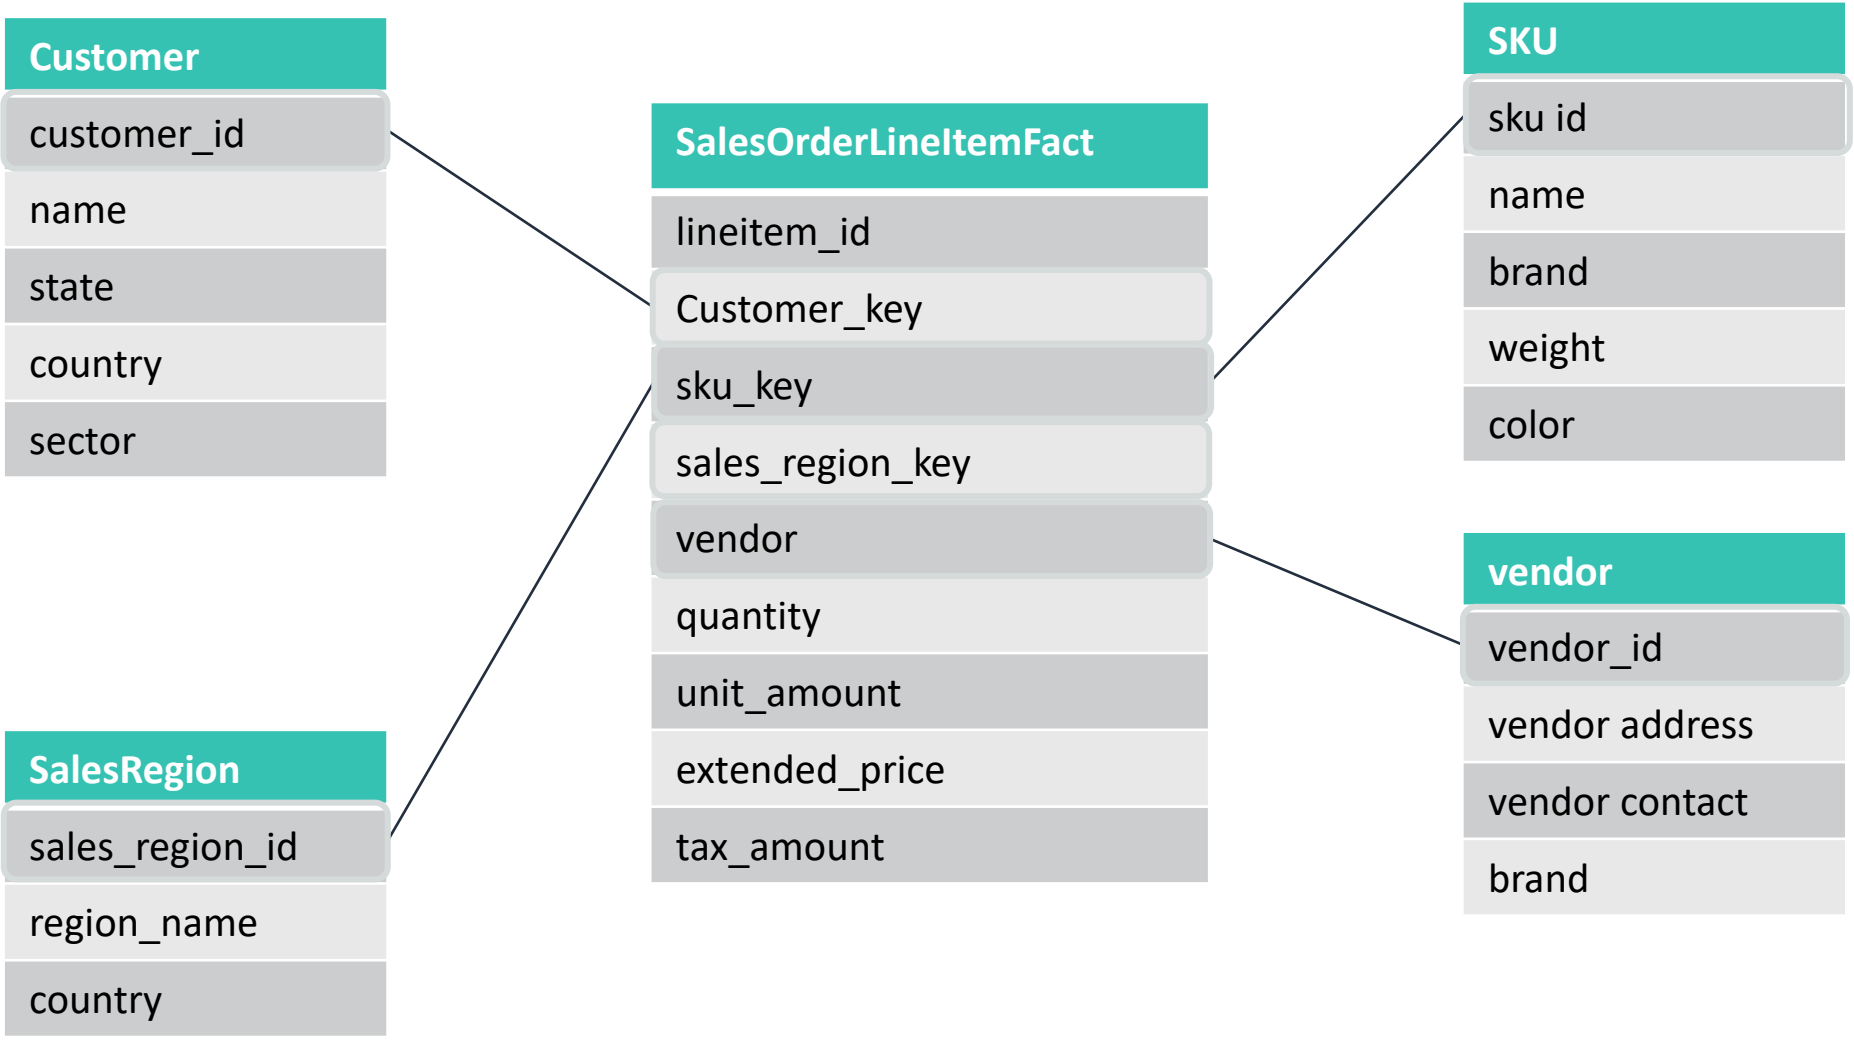

Normalized Modeling: Snowflake Schema

- Developed by “the inventor of the relation model,” Edgar F. Codd.
- Data is stored in related tables that are normalized to third normal form (3NF) to reduce data redundancy.

Snowflake Schema Example

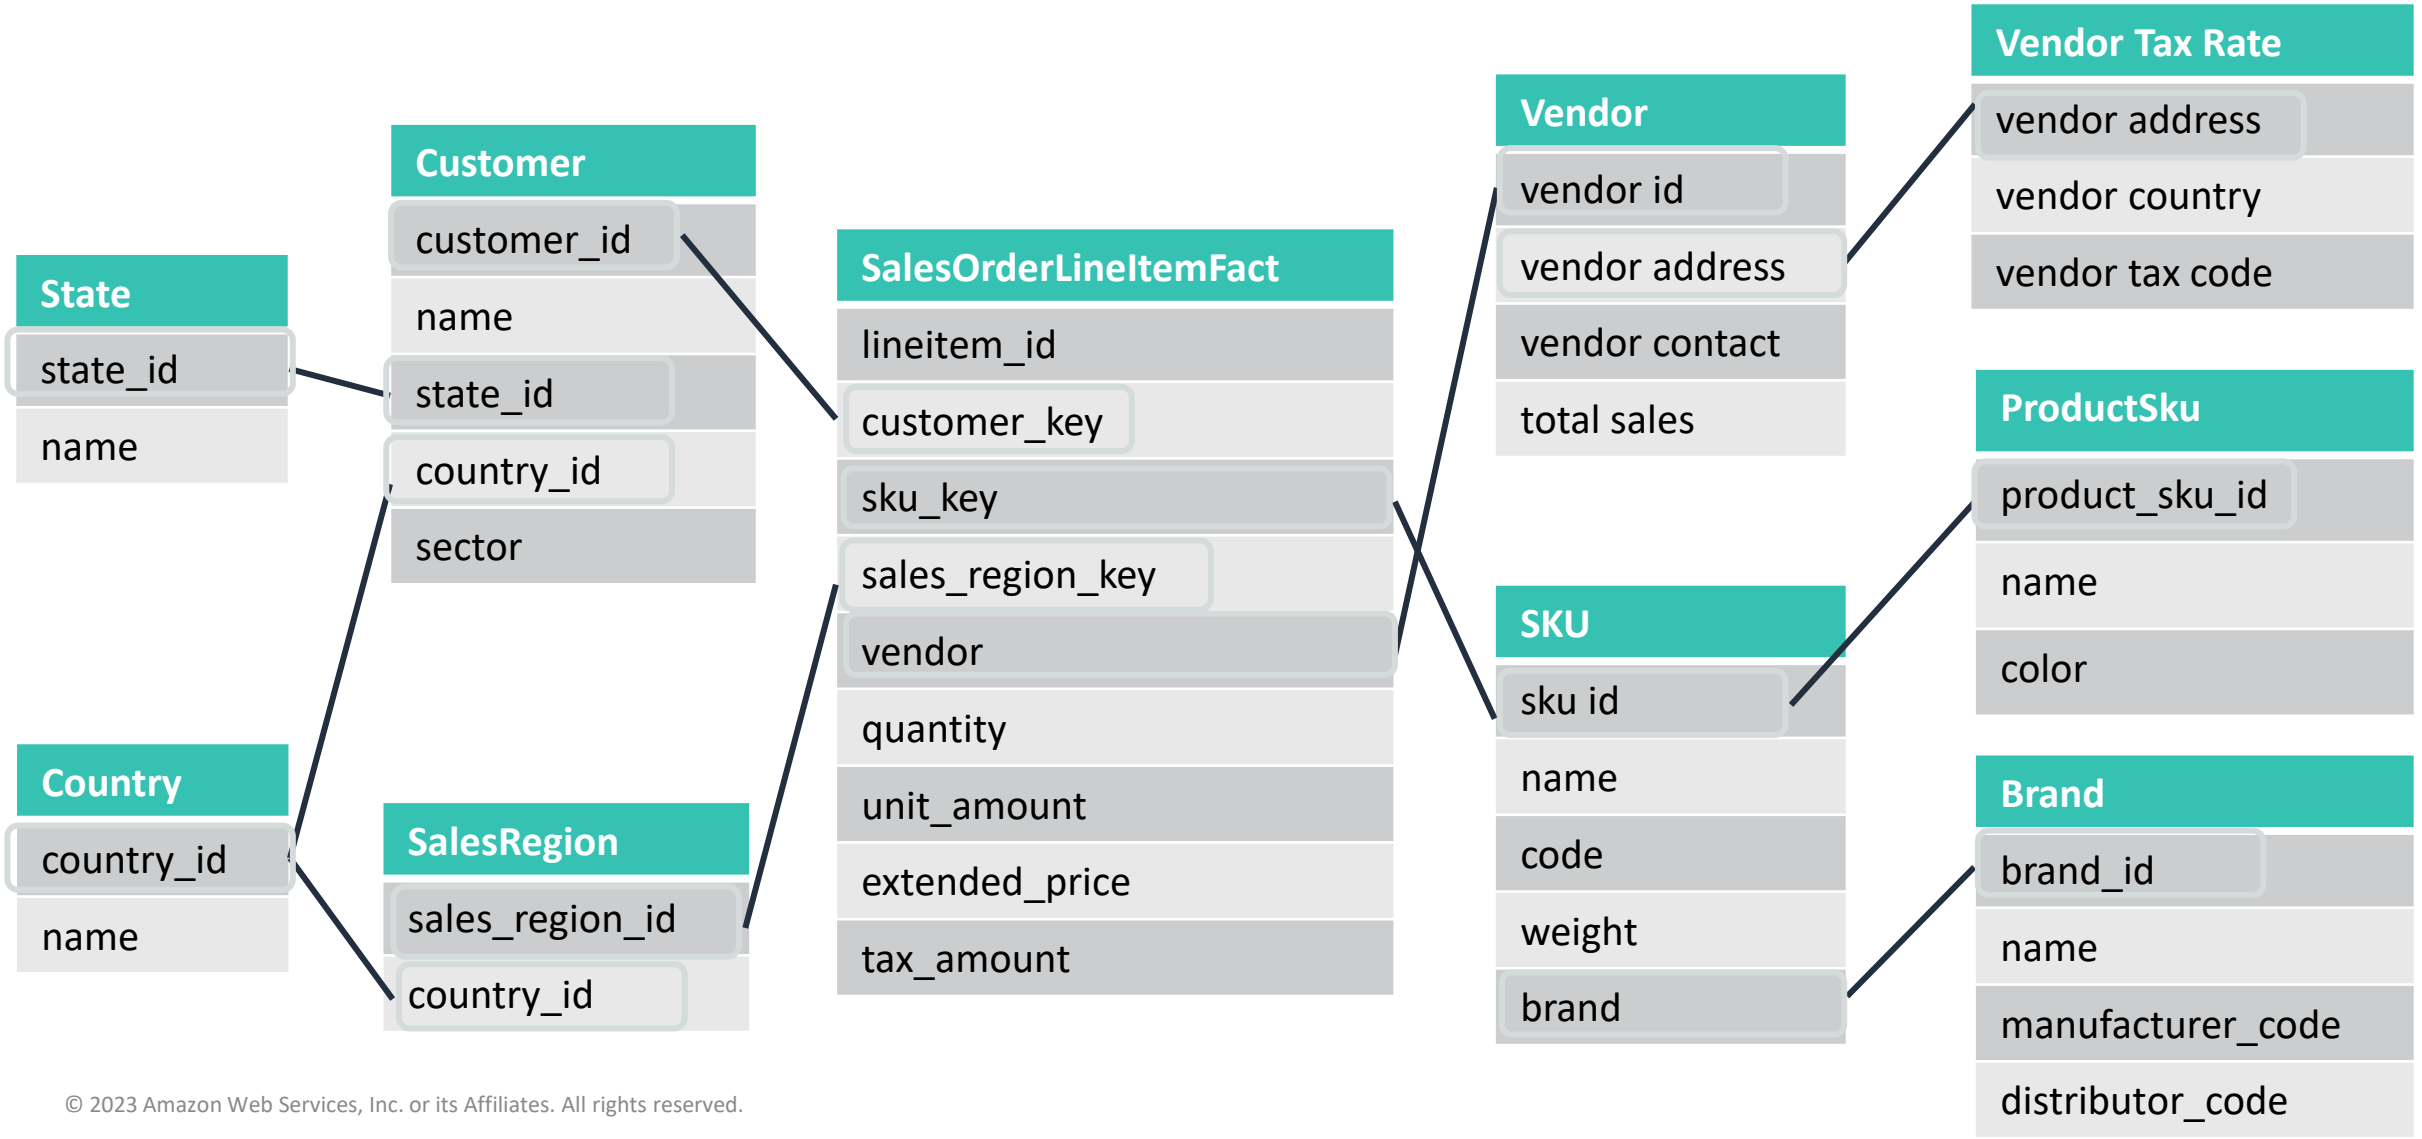

Data Models and Redshift

aws training and certification

- Data models were created to handle the complexity of organizing data from multiple data sources.
- Amazon Redshift already works with these common schemas and more. We recommend that you load your data in Amazon Redshift to try it out and adjust as needed.

The Intersection of Data Warehousing and Big Data

aws training and certification

What Is Big Data?

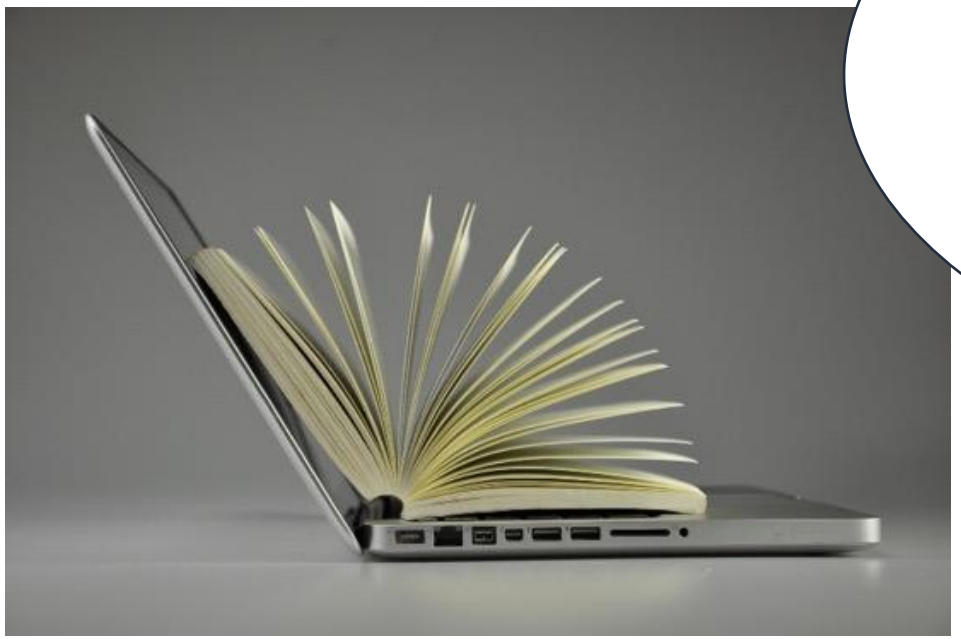

"Big data is high-volume, high-velocity, and/or high-variety information assets that demand cost-effective, innovative forms of information processing that enable enhanced insight, decision making, and process automation."

Gartner IT Glossary

Unconstrained Big Data Growth

- IT/Application server logs
  - IT Infrastructure logs, Metering, Audit logs, Change logs
- Websites/Mobile apps/Ads
  - Clickstream, User Engagement
- Sensor data/IoT
  - Weather, Smart Grids, Wearables
- Social media, user content
  - 450MM+ Tweets/day

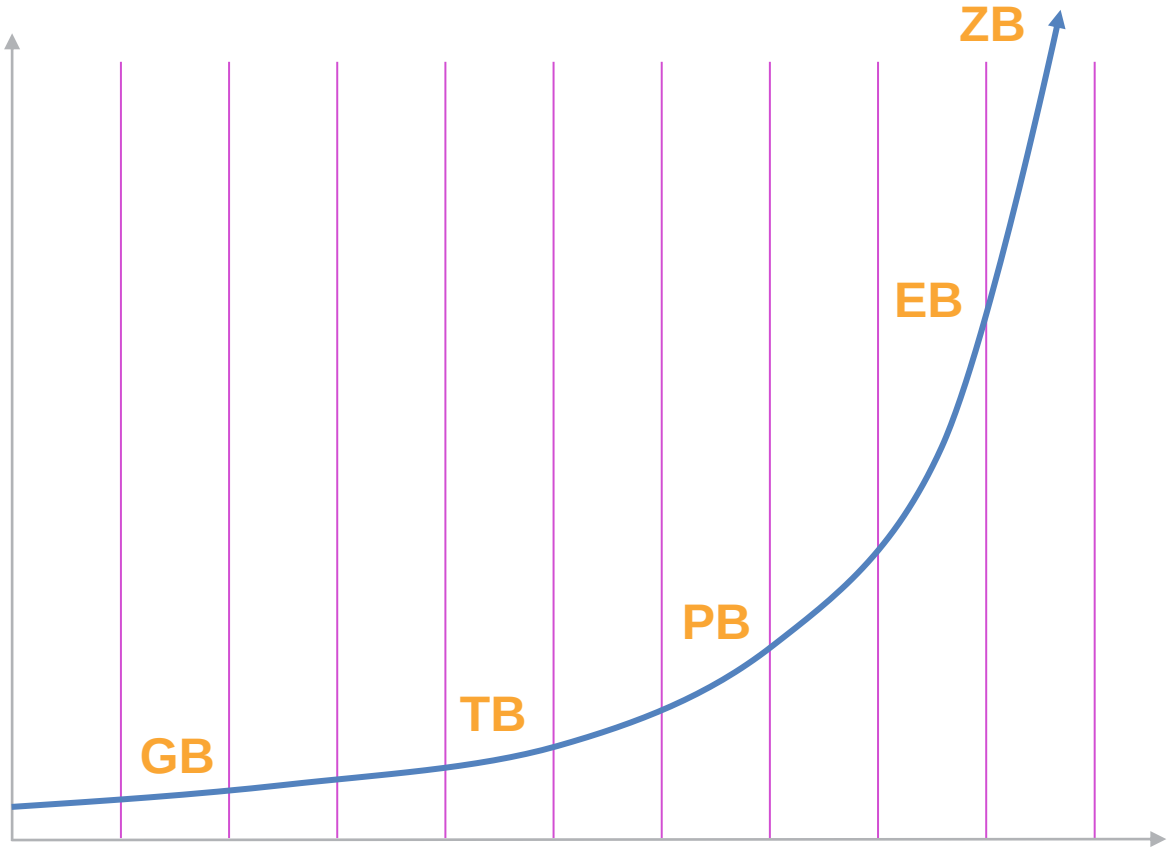

Characteristics of Big Data

aws training and certification

The 5 Vs

- **V**olume
- **V**ariety
- **V**elocity
- **V**ariability
- **V**eracity

Typical Use Cases

- Customer segmentation
- Marketing spend optimization
- Financial modeling and forecasting
- Ad targeting and real-time bidding
- Clickstream analysis
- Fraud detection

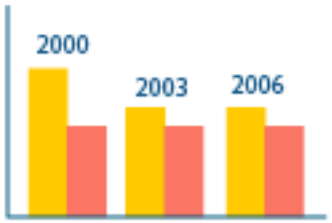

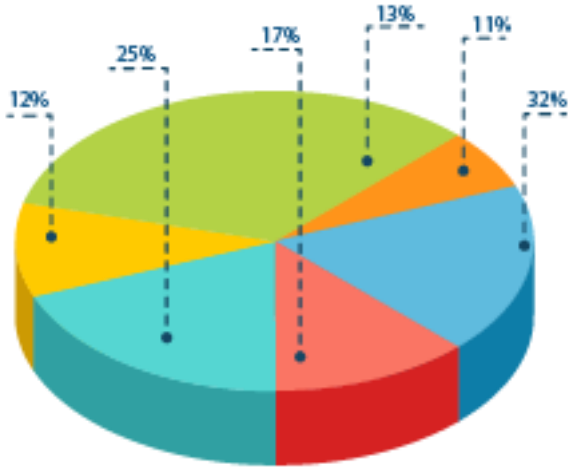

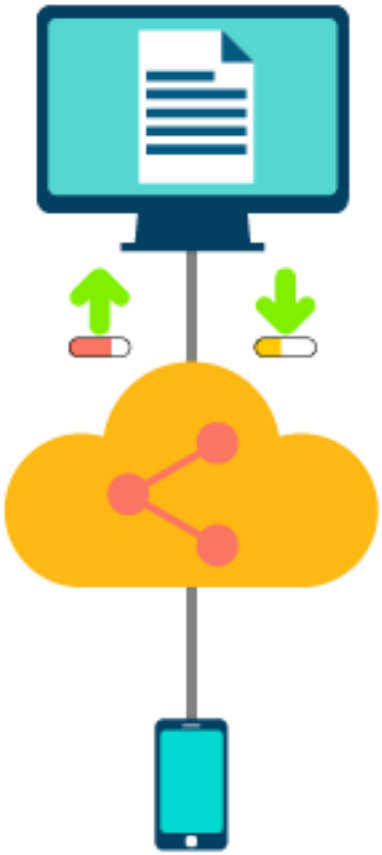

Metrics

- Visits, views, clicks, and purchases
- Sources, devices, locations, times
- Latency, throughput, uptime
- Likes, shares, friends, follows
- Prices, frequency

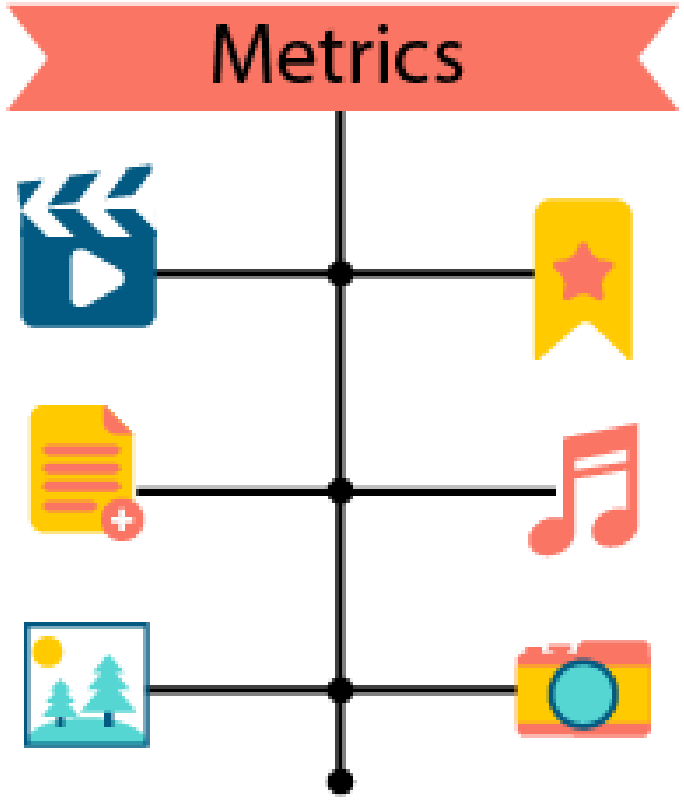

Data Sources

- Relational databases
- NoSQL databases
- Web servers
- Mobile phones
- Tablets
- Data feeds

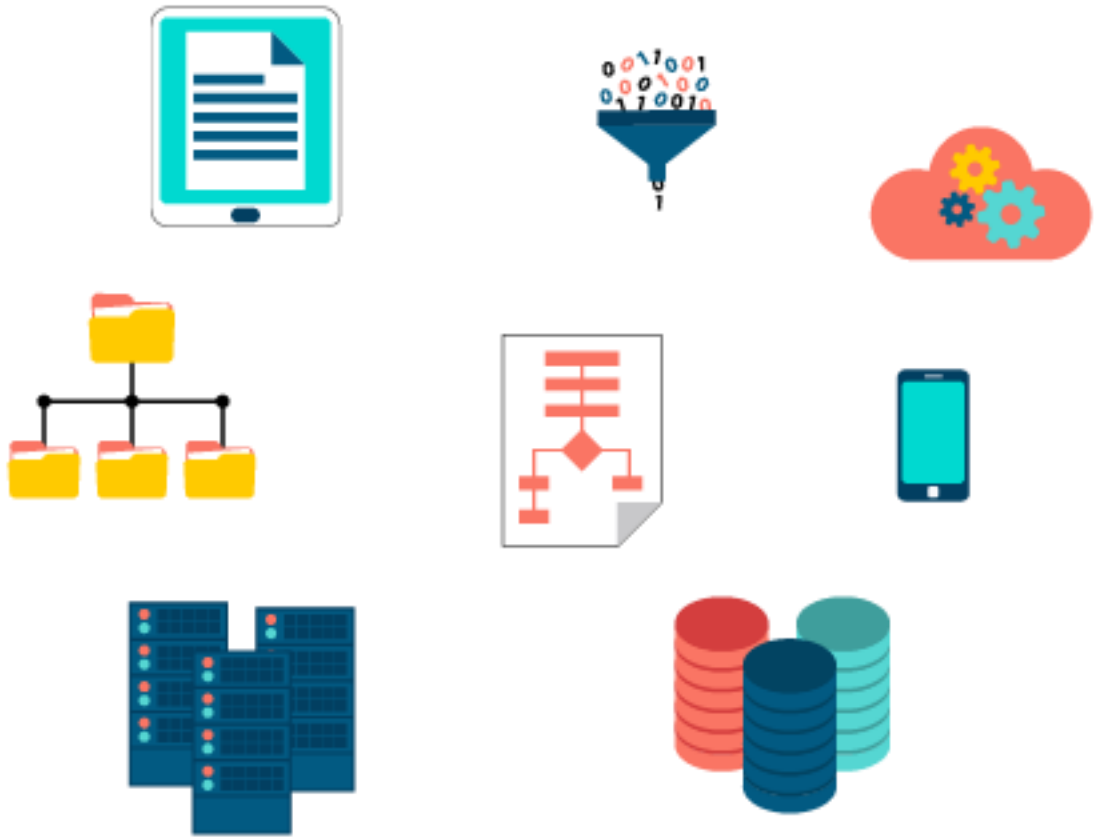

Data Formats

- Structured, semi-structured, and unstructured
- Text
- Binary
- Streaming and near real-time
- Batched

Use the Right Tool for the Job

Compute-intensive

Streaming analytics

Data transformation

Multi-input operations

Big Data vs Data Warehouses

Big Data is a concept.

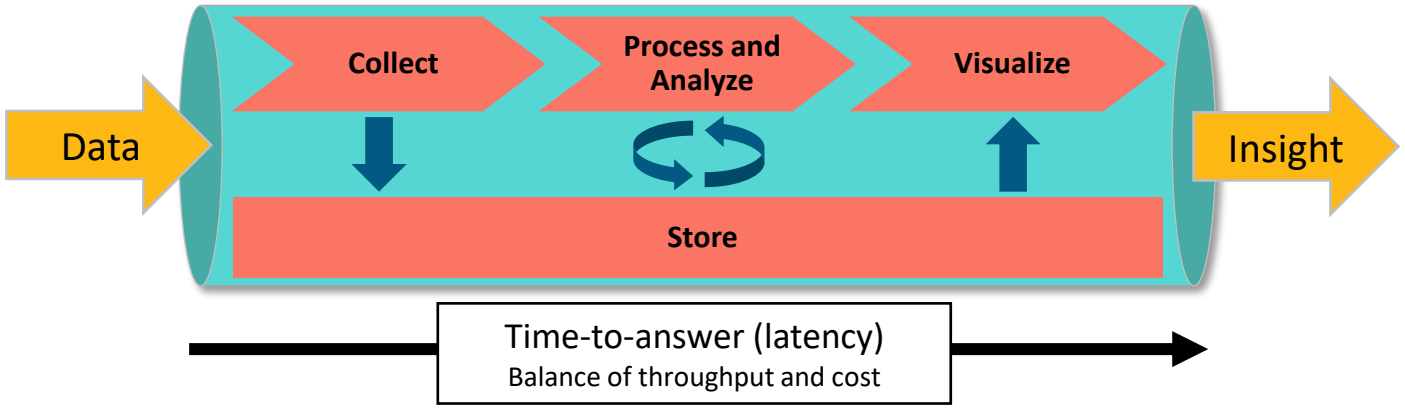

A data warehouse:

- Can be used with both small and large datasets
- Can be used in a Big Data system

Overview of Data Management in AWS

aws training and certification

AWS Global Infrastructure: Current Regions

aws training and certification

AWS Global Infrastructure: Edge Locations

aws training and certification

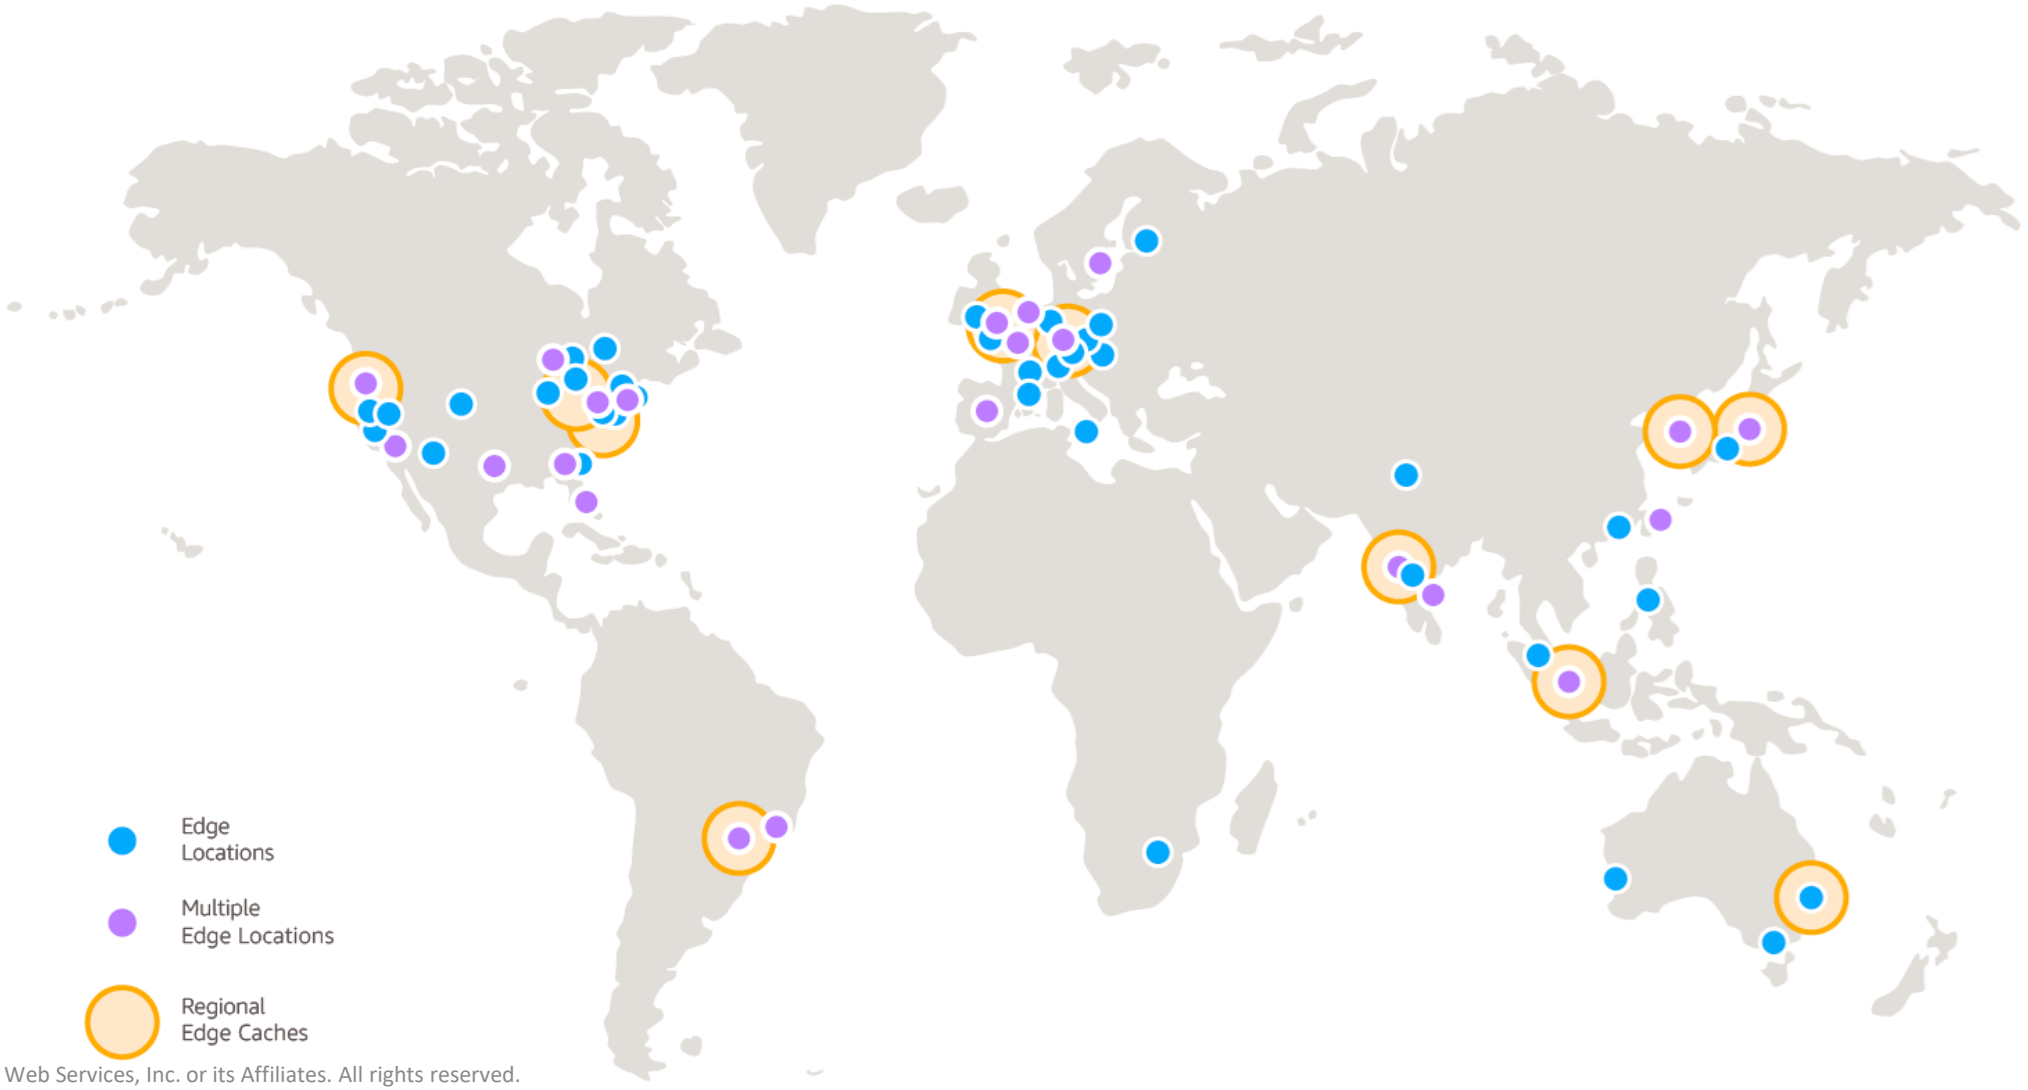

On-Premises Systems

aws training and certification

In traditional systems, you are responsible for the following tasks:

- Application Development and Optimization
- OS patching
- Scaling
- OS installation
- High availability
- Server maintenance
- Database backups
- Rack and stack
- DB software patches
- Power, HVAC, Network
- DB software installs

Moving Systems to AWS

aws training and certification

By redeploying your systems in Amazon EC2 instances, you unload some of the tasks to AWS.

Your Responsibilities

- Application Development and Optimization
- OS patching
- Scaling
- High availability
- Database backups
- DB software patches
- DB software installs

AWS Responsibilities

- OS installation
- Server maintenance
- Rack and stack
- Power, HVAC, Network

AWS Managed Services

aws training and certification

By transitioning your systems to AWS-managed services, you unload many tasks to AWS and can focus on your applications and optimizations for your business.

Your Responsibilities

Application Development and Optimization

AWS Responsibilities

- OS patching
- OS installation
- Server maintenance
- Rack and stack
- Power, HVAC, Network
- Scaling
- High availability
- Database backups
- DB software patches
- DB software installs

Benefits of AWS Database Services

aws training and certification

Managed

AWS installs, patches, and manages the services.

Scalable

You can grow or shrink resources as needed.

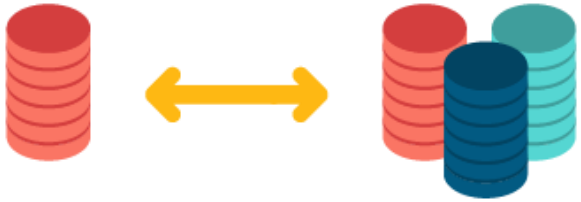

No up-front cost

You pay only for what you use.

Integrated

The database services are already integrated with other AWS services.

Traditional Data Flow Overview

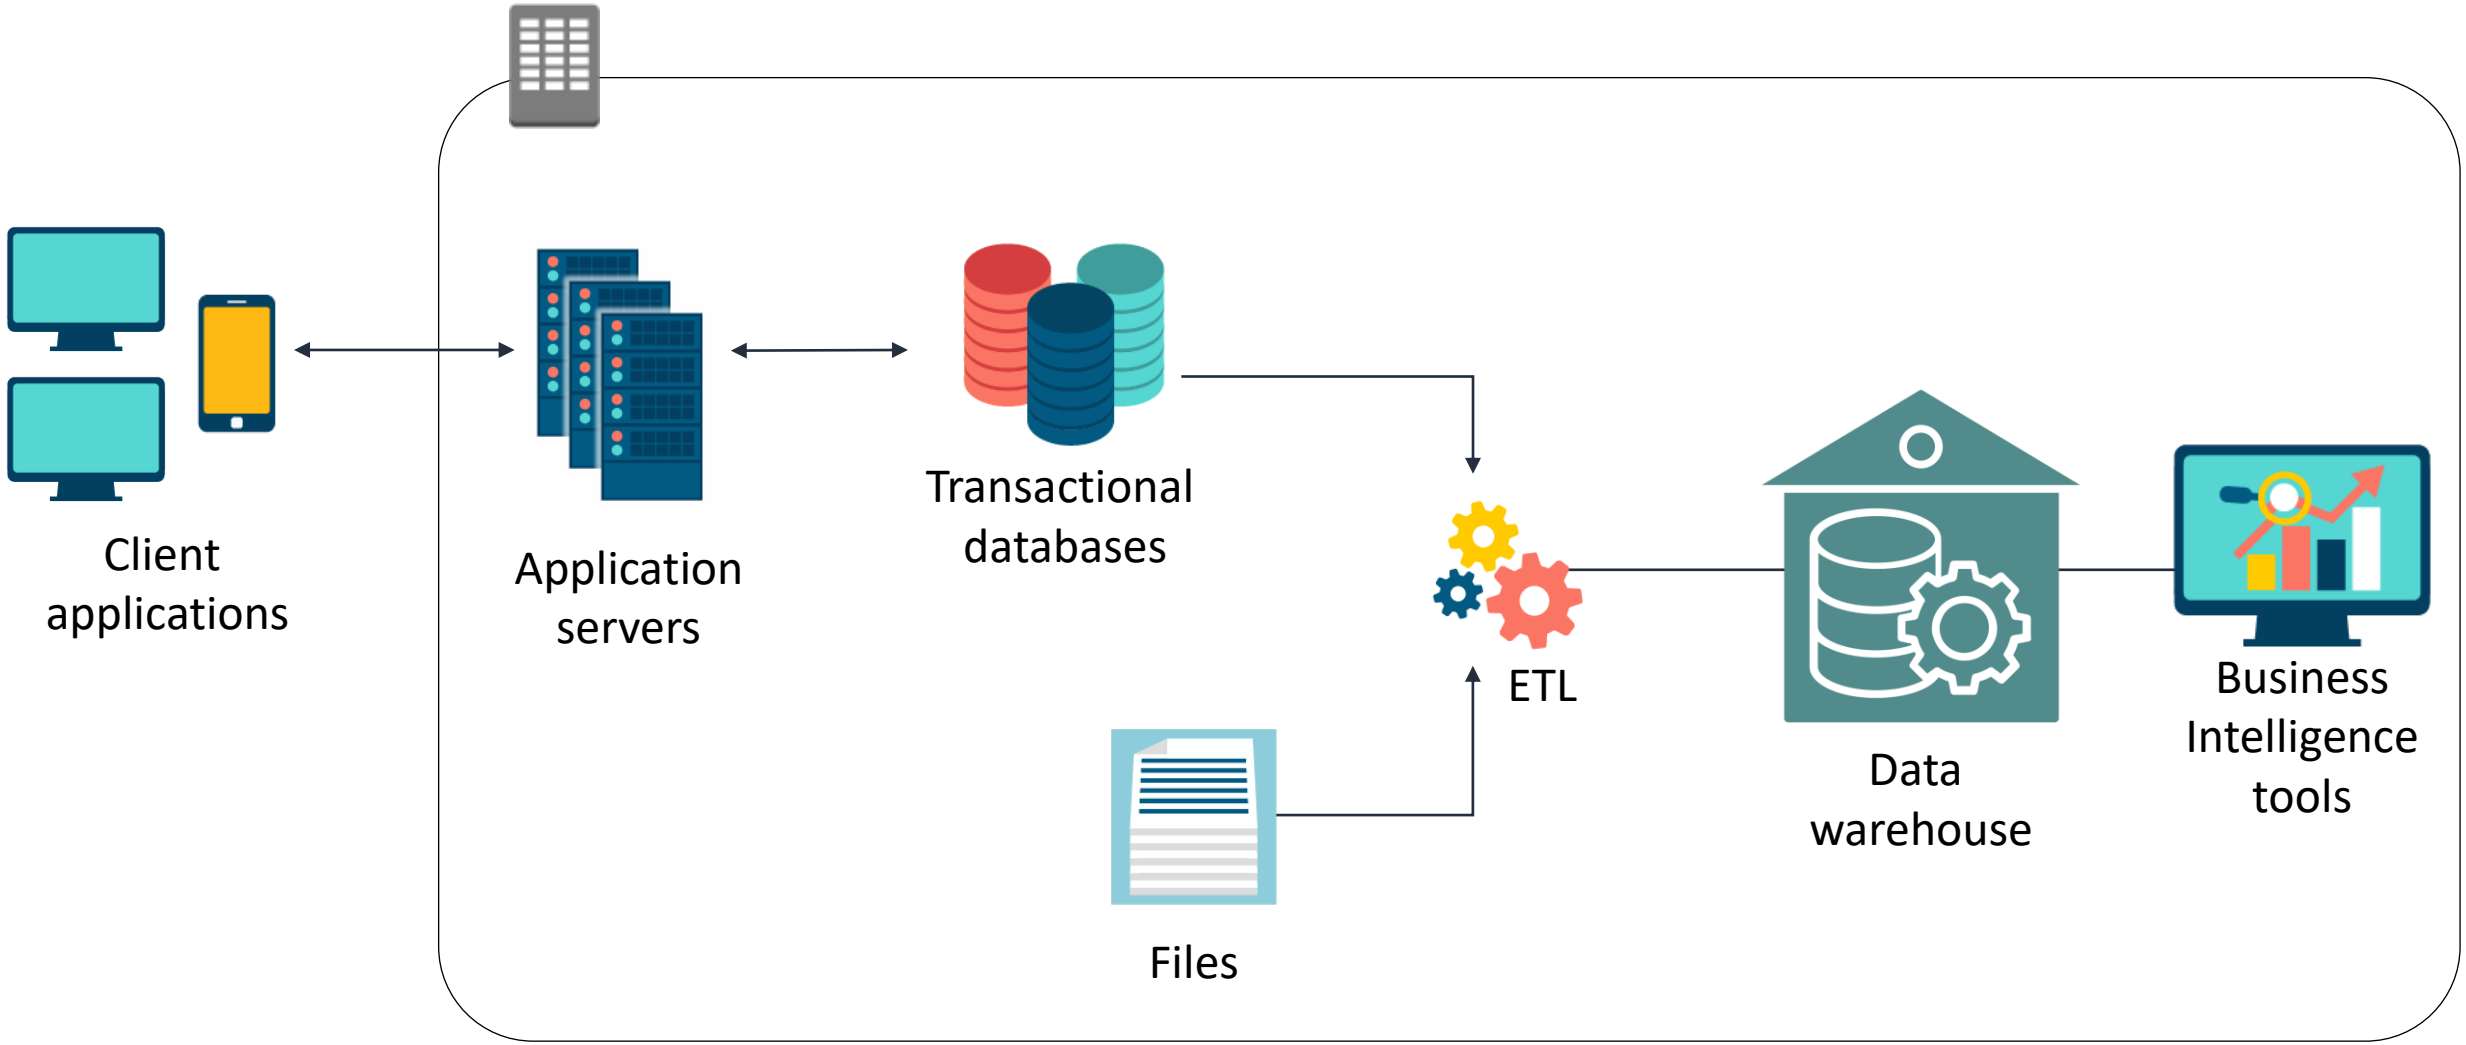

Data Flow With Traditional Systems

aws training and certification

Collect

Application servers

Store

Transactional databases

Files

Process

ETL

Analyze

Business Intelligence tools

Data warehouse

Data Flow With AWS Services

aws training and certification

Collect

Amazon Kinesis

AWS Direct Connect

AWS Snowball

Store

Amazon S3

Amazon Glacier

Amazon DynamoDB

Amazon RDS

SSH-Enabled Host

Process

AWS Lambda

Amazon EMR

AWS Data Pipeline

Amazon Athena

AWS Glue

Analyze

Amazon Redshift

Amazon EMR

Amazon Machine Learning

Amazon QuickSight

RedShift

aws training and certification

Amazon Redshift Is a…

aws training and certification

fast,

simple,

cost-effective,

fully managed,

petabyte-scale,

enterprise-grade,

relational

…data warehousing service on AWS.

Amazon Redshift

Data Sources Integrated with Redshift

aws training and certification

These services are integrated data sources loading data in parallel into Amazon Redshift. An SSH-enabled host can be either an Amazon EC2 instance or an on-premises computer.

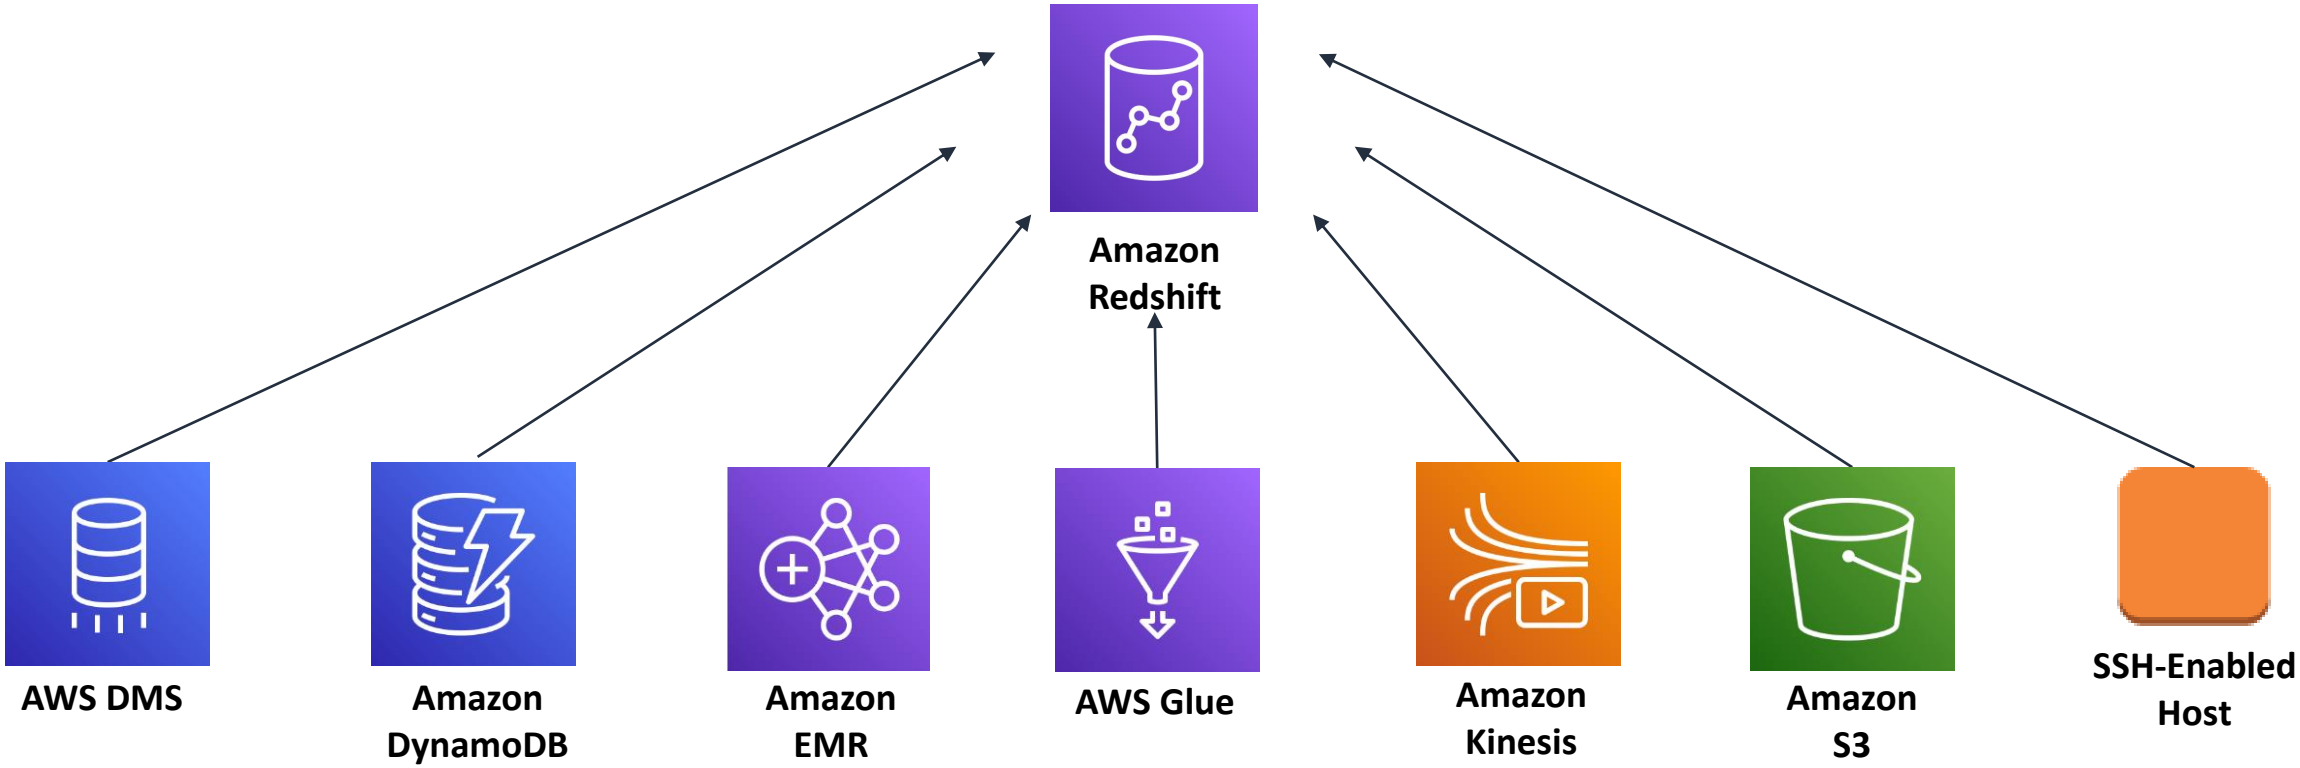

Interacting with Amazon Redshift

aws training and certification

Management Interfaces

AWS Management Console

AWS CLI

AWS SDKs

Amazon Redshift Query API

Query and Analysis Tools

SQL tools

Business Intelligence Applications

Amazon QuickSight

Massively Parallel Processing (MPP)

aws training and certification

Amazon Redshift is built on MPP architecture, which:

- Is a distributed, shared-nothing architecture that is easily scalable by adding or removing nodes
- Is optimized for analytic workloads
- Provides the ability to distribute, scan, and process queries in parallel across nodes in the cluster
- Improves query performance significantly for complex analytical queries against massive data sets

Cluster Overview

An Amazon Redshift cluster comprises one leader node, and one or more compute nodes.

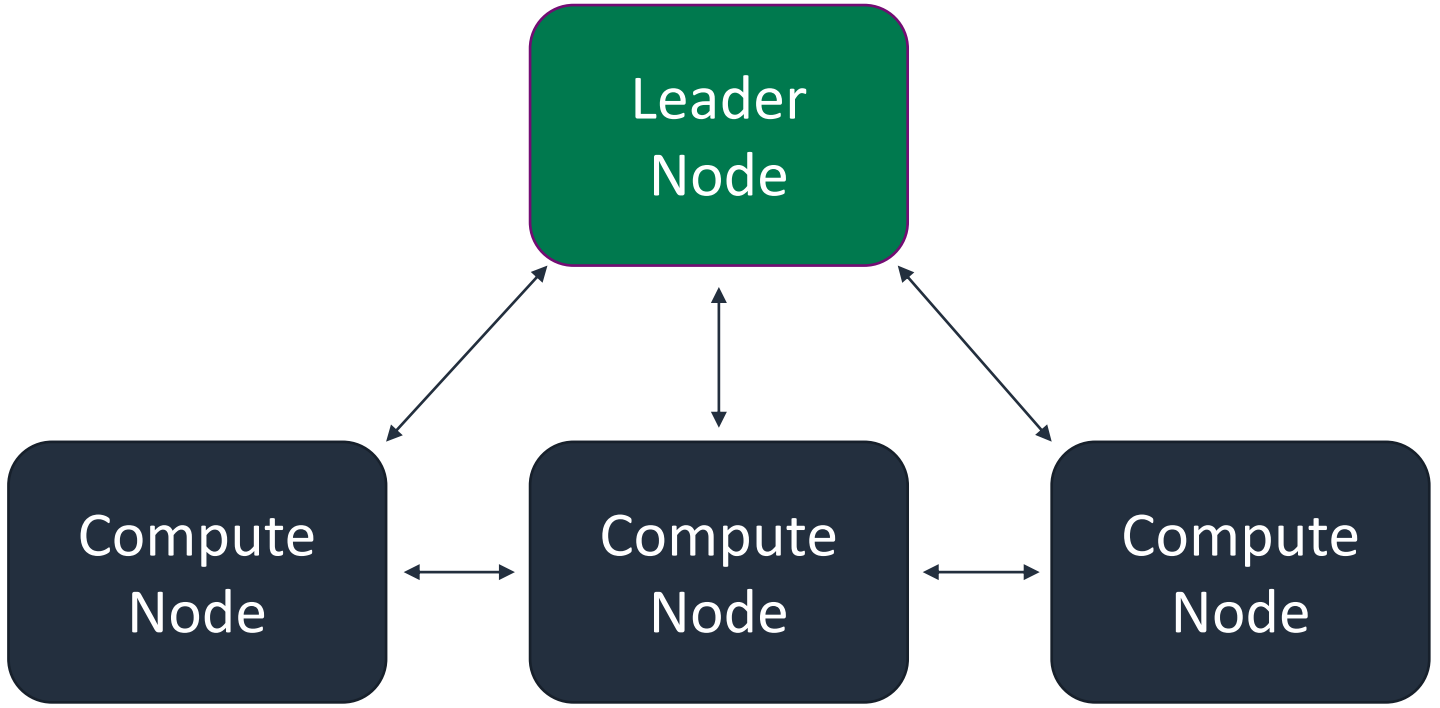

Node Slices

A compute node is partitioned into slices.

Each slice is allocated a portion of the cluster's memory and disk space.

The Leader node distributes data to the slices.

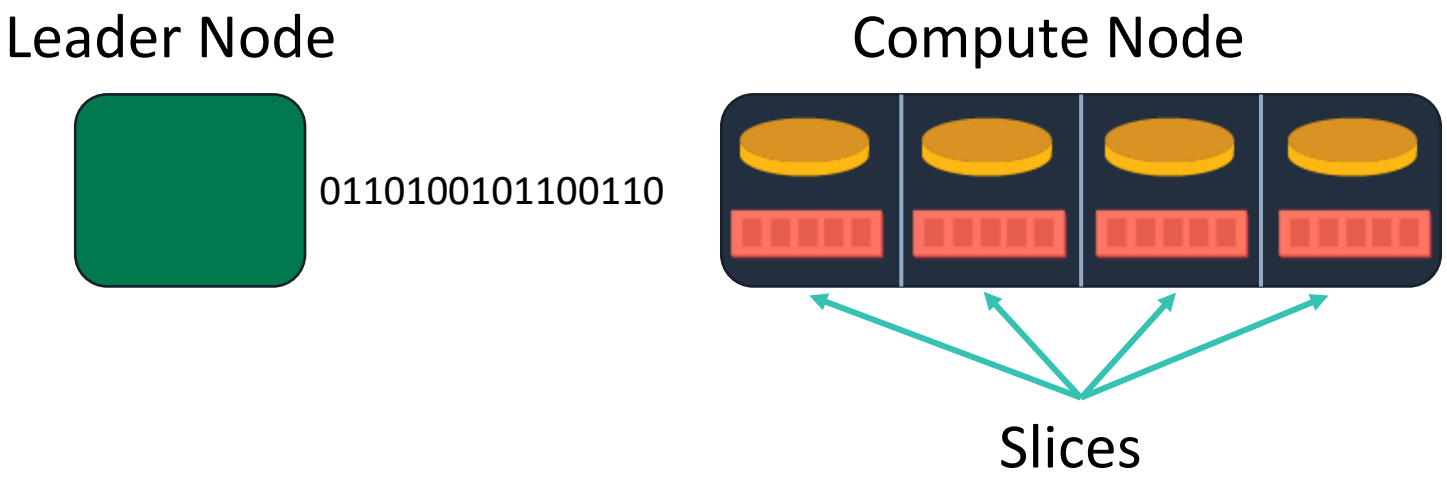

Query Processing

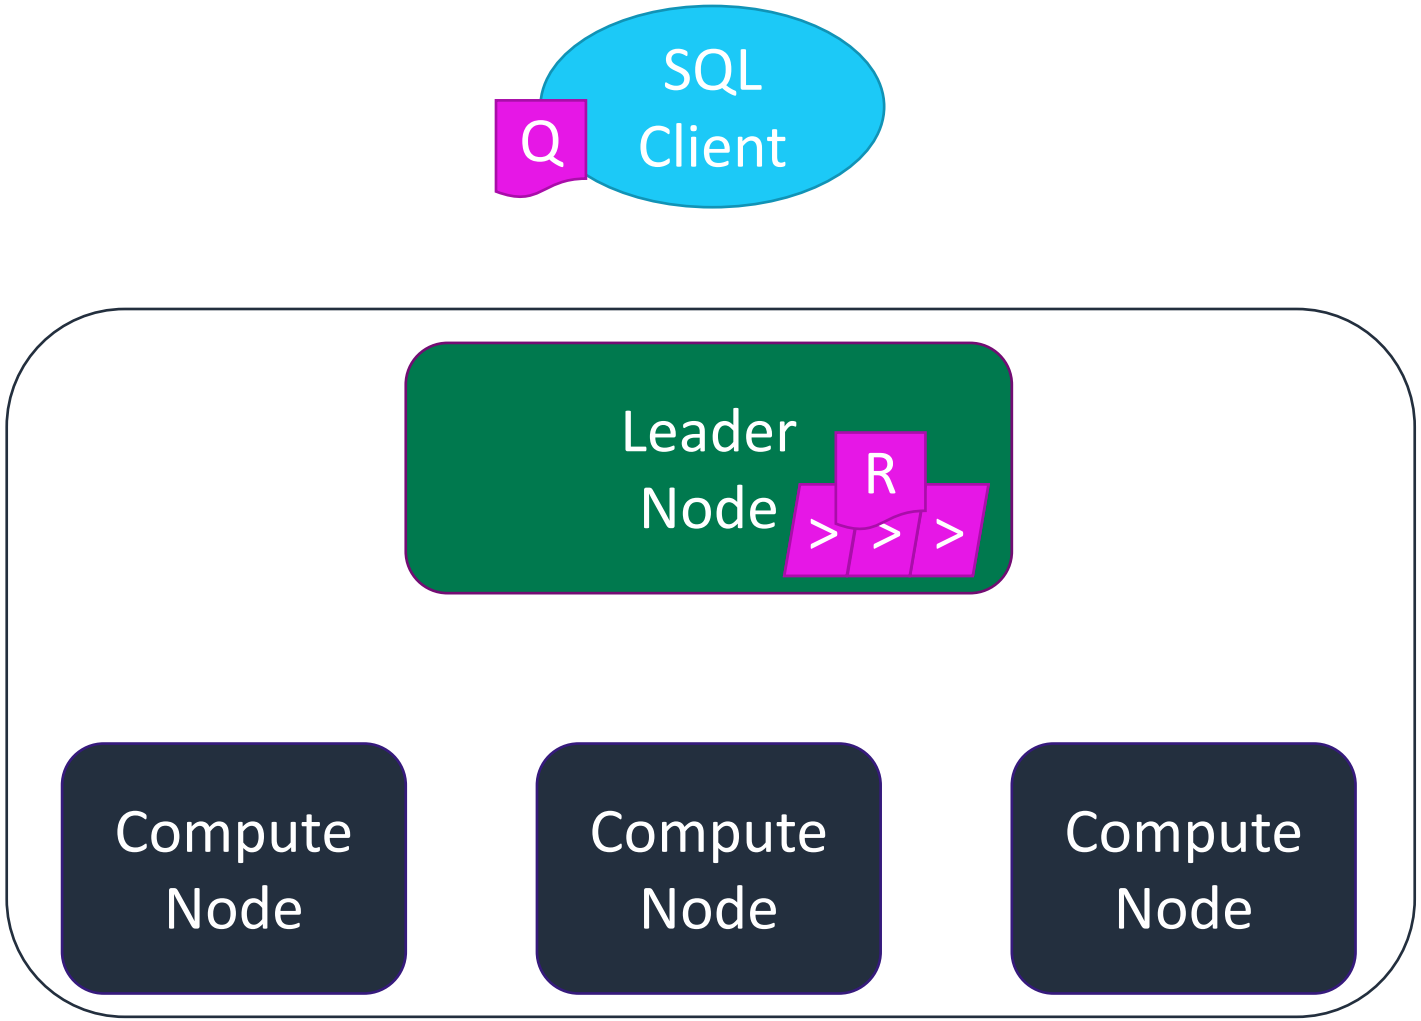

1. The SQL client submits a query.
2. The leader node parses the query and develops an execution plan.
3. The leader node distributes work among compute nodes.
4. The compute nodes run the query according to the plan, and return interim results to the leader node.
5. The leader node aggregates and returns results to the SQL client.

Internal Architecture

Network

- SQL endpoint
- No direct client application communication
- Client communication must be JDBC/ODBC
- Can be encrypted

Leader Node

- Execution engine
- Stores metadata
- Coordinates query execution

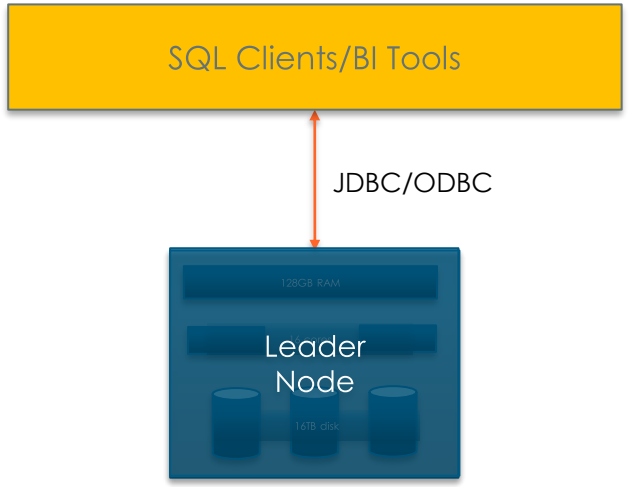

Internal Architecture

Compute nodes

- Local, columnar storage
- Run queries in parallel
- Load, backup, restore via Amazon S3
- Load from Amazon DynamoDB or SSH
- Fault tolerant

Three hardware platforms

- Optimized for data processing
- DS2: HDD; scale from 2 TB to 2 PB
- DC2: SSD; scale from 160 GB to 326 TB
- RA3: SSD; scale from 64 TB to 16,384 TB (16.38 PB)

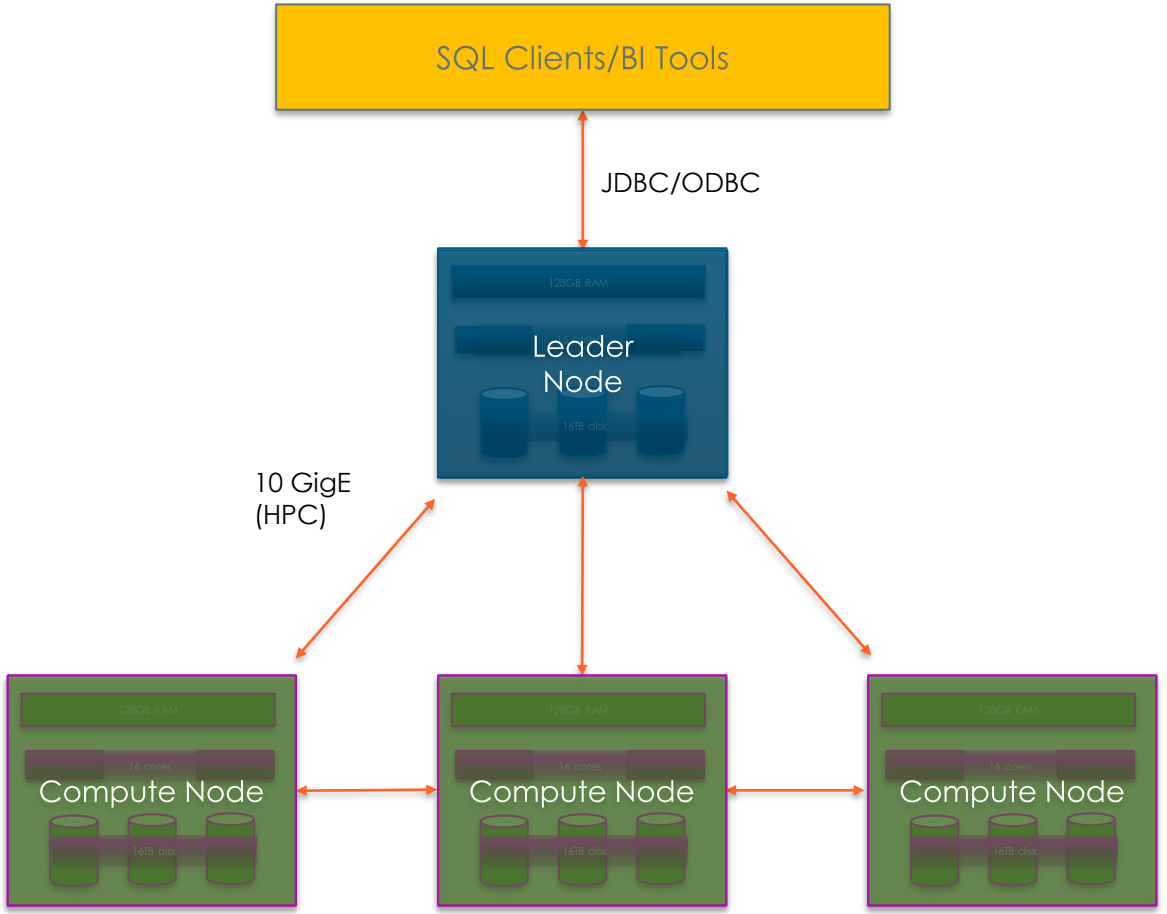

Columnar Storage

- Is optimized for scanning large data sets and complex analytic queries
- Enables a data block to store and compress significantly more values for a column compared to row-based storage
- Eliminates the need to read redundant data by reading only the columns that you include in your query
- Offers overall performance benefits that can help eliminate the need to aggregate data into cubes as in some other OLAP systems

Columnar Compression

aws training and certification

Amazon Redshift supports the following compression encodings:

- Raw
- Byte-Dictionary
- Delta
- LZO
- Mostly
- Runlength
- Text
- Zstandard
- AZ64 (A new compression algorithm to save disk space)

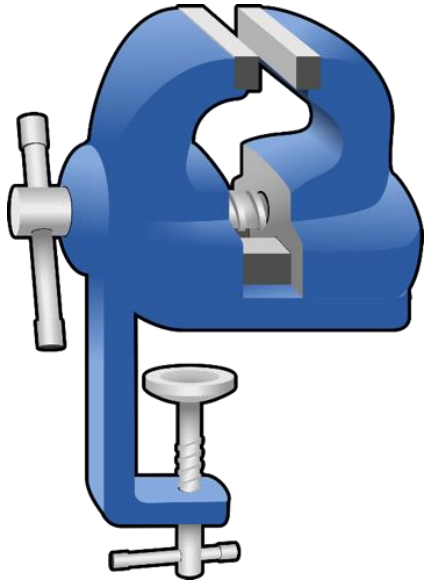

Analyzing Near Real–Time Tweets

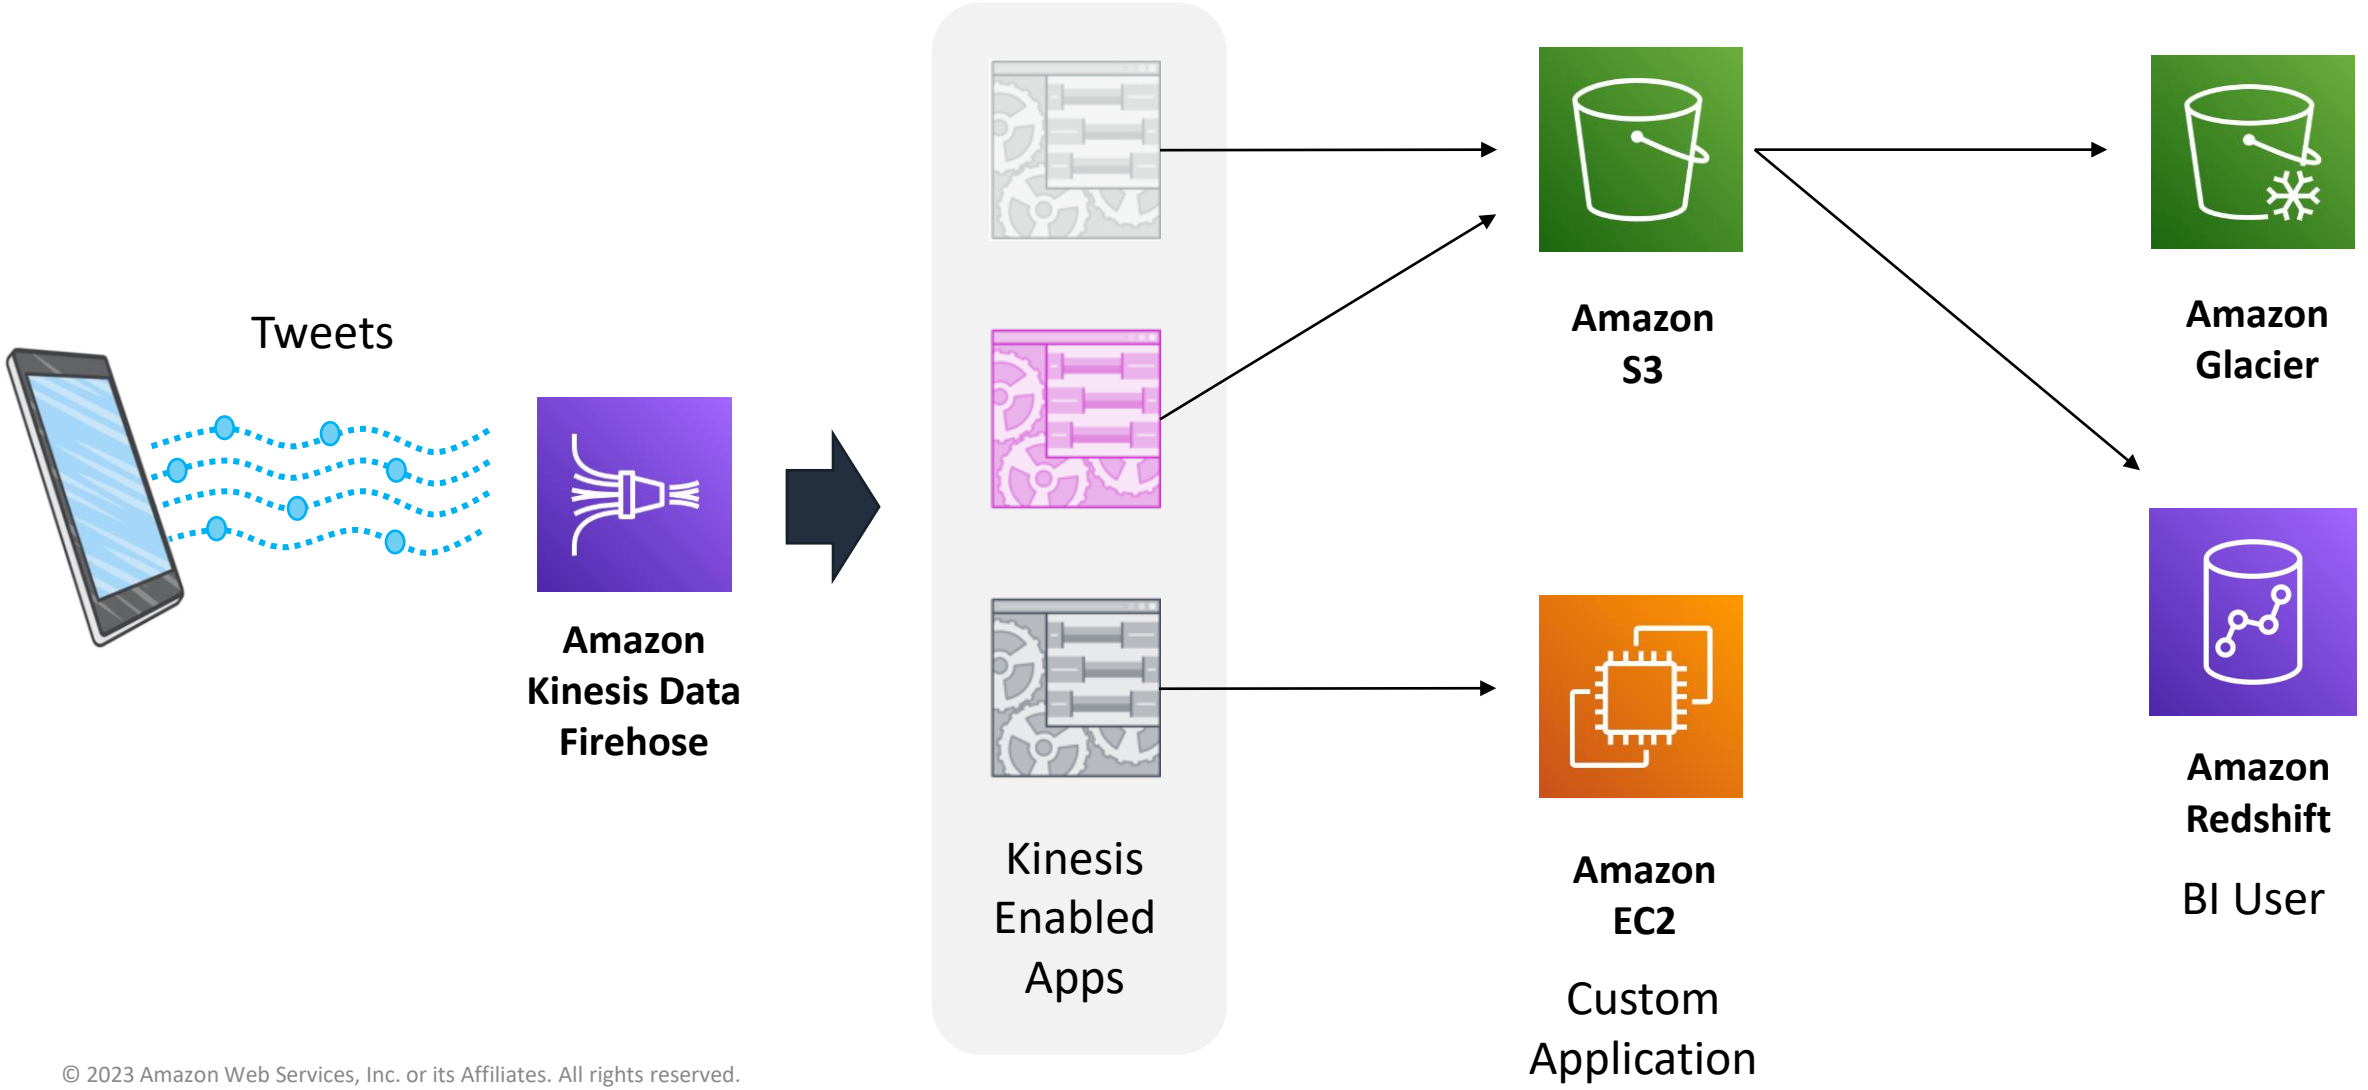

Analyzing Near Real–Time Tweets from Twitter Firehose

aws training and certification

A look at the numbers:

- 500 million tweets per day = ~5,800 tweets per second
- 2000 tweets = ~12 MB per second (~1 TB per day)
- $0.015 per hour per shard
- $0.028 per million PUT requests
- $0.765 per hour for Amazon Kinesis
- $0.850 per hour for a 2 TB node in Amazon Redshift
- $1.28 per hour for uncompressed data in Amazon S3

Total: $2.895 per hour